



Question 3

```
# @title Install & Imports
!pip -q install torch torchvision torchaudio --index-url https://download.pyt
!pip -q install scikit-learn matplotlib tqdm >/dev/null

import os, time, random, math
import numpy as np
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader, Subset, TensorDataset
from torchvision import datasets, transforms, models
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix, Confus
from sklearn.manifold import TSNE
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
import matplotlib.pyplot as plt
from tqdm.auto import tqdm

# Reproducibility
SEED = 42
random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed_all(SEED)

DEVICE = "cuda" if torch.cuda.is_available() else "cpu"
print("Device:", DEVICE)

# Output directory
OUT_DIR = "/content/outputs"
os.makedirs(OUT_DIR, exist_ok=True)

# Small helpers
def count_params(model, only_trainable=False):
    params = [p.numel() for p in model.parameters() if (p.requires_grad or no
    return sum(params))

def time_inference(model, loader, device=DEVICE):
    model.eval()
    torch.cuda.synchronize() if device=="cuda" else None
    t0 = time.time()
    with torch.no_grad():
        for xb, yb in loader:
            xb = xb.to(device)
            _ = model(xb)
```

```

torch.cuda.synchronize() if device=="cuda" else None
return time.time()-t0

def plot_confmat(y_true, y_pred, title, fname):
    cm = confusion_matrix(y_true, y_pred, labels=list(range(10)))
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=list(range(10)))
    fig, ax = plt.subplots(figsize=(6,6))
    disp.plot(ax=ax, colorbar=False)
    ax.set_title(title)
    plt.tight_layout()
    fig.savefig(fname, dpi=150)
    plt.show()

def tsne_plot(z, labels, title, fname):
    z2 = TSNE(n_components=2, init="pca", random_state=SEED, perplexity=30).fit_transform(z)
    plt.figure(figsize=(6,5))
    scatter = plt.scatter(z2[:,0], z2[:,1], c=labels, s=6, alpha=0.8)
    plt.title(title)
    plt.xticks([]); plt.yticks([])
    plt.tight_layout()
    plt.savefig(fname, dpi=150)
    plt.show()

Device: cpu

# @title Download MNIST & Fashion-MNIST
BATCH = 128

tfm_basic = transforms.Compose([
    transforms.ToTensor(), # [0,1]
    transforms.Normalize((0.1307,), (0.3081,)) # standard MNIST stats
])

mnist_train = datasets.MNIST(root="/content", train=True, download=True, transform=tfm_basic)
mnist_test = datasets.MNIST(root="/content", train=False, download=True, transform=tfm_basic)

fmnist_test = datasets.FashionMNIST(root="/content", train=False, download=True, transform=tfm_basic)

# Optional: stratified subset of training to speed up (set to None to use full)
SUBSET_PER_CLASS = 1000 # e.g., 1000 => 10k total. Set to None to use full.
if SUBSET_PER_CLASS is not None:
    targets = np.array(mnist_train.targets)
    idxs = []
    for c in range(10):
        c_idx = np.where(targets==c)[0]
        np.random.shuffle(c_idx)
        idxs.extend(c_idx[:SUBSET_PER_CLASS])
    mnist_train = Subset(mnist_train, idxs)
    print(f"Using stratified subset: {len(mnist_train)} samples")

```

```

else:
    print("Using full training set")

train_loader = DataLoader(mnist_train, batch_size=BATCH, shuffle=True, num_workers=4)
test_loader = DataLoader(mnist_test, batch_size=BATCH, shuffle=False, num_workers=4)

# Also create flat versions (for sklearn baselines)
def dataset_to_numpy(dset):
    xs, ys = [], []
    for x, y in DataLoader(dset, batch_size=512, shuffle=False):
        xs.append(x.numpy())
        ys.append(y.numpy())
    X = np.concatenate(xs, axis=0) # N,1,28,28
    y = np.concatenate(ys, axis=0)
    return X, y

Xtr, ytr = dataset_to_numpy(mnist_train)
Xte, yte = dataset_to_numpy(mnist_test)
Xtr_flat = Xtr.reshape(len(Xtr), -1)
Xte_flat = Xte.reshape(len(Xte), -1)
print("Train flat:", Xtr_flat.shape, "Test flat:", Xte_flat.shape)

# Fashion-MNIST test in numpy (for cross-domain eval with MLP embeddings)
Xf_te, yf_te = dataset_to_numpy(fmnist_test)

Using stratified subset: 10000 samples
Train flat: (10000, 784) Test flat: (10000, 784)

```

3.1

```

# @title Define MLP (30→20→10) and trainer
class MLP(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()
        self.fc1 = nn.Linear(28*28, 30)
        self.fc2 = nn.Linear(30, 20)
        self.fc3 = nn.Linear(20, 10)
    def features20(self, x): # returns the 20-dim hidden layer
        x = self.flatten(x)
        x = F.relu(self.fc1(x))
        x = self.fc2(x) # no ReLU here so we can see raw embeddings
        return x
    def forward(self, x):
        z = self.features20(x)
        x = F.relu(z)
        x = self.fc3(x)
        return x

```

```

def train_classifier(model, loader, epochs=5, lr=1e-3, wd=0.0, device=DEVICE):
    model.to(device)
    opt = torch.optim.Adam(model.parameters(), lr=lr, weight_decay=wd)
    criterion = nn.CrossEntropyLoss()
    model.train()
    for ep in range(epochs):
        pbar = tqdm(loader, leave=False)
        running = 0.0
        for xb, yb in pbar:
            xb, yb = xb.to(device), yb.to(device)
            opt.zero_grad()
            logits = model(xb)
            loss = criterion(logits, yb)
            loss.backward()
            opt.step()
            running += loss.item()*xb.size(0)
            pbar.set_description(f"Epoch {ep+1}/{epochs} loss={running/((pbar.n - pbar.n) + 1)}")
    return model

def predict(model, loader, device=DEVICE):
    model.eval()
    preds, golds = [], []
    with torch.no_grad():
        for xb, yb in loader:
            logits = model(xb.to(device))
            pred = logits.argmax(dim=1).cpu().numpy()
            preds.append(pred)
            golds.append(yb.numpy())
    return np.concatenate(preds), np.concatenate(golds)

# @title Train MLP and evaluate against RandomForest & LogisticRegression
# MLP
mlp = MLP()
mlp = train_classifier(mlp, train_loader, epochs=10, lr=1e-3, wd=1e-4)

y_pred_mlp, y_true = predict(mlp, test_loader)
acc_mlp = accuracy_score(y_true, y_pred_mlp)
f1_mlp = f1_score(y_true, y_pred_mlp, average="macro")
print(f"MLP - Acc: {acc_mlp:.4f}, F1(macro): {f1_mlp:.4f}")
plot_confmat(y_true, y_pred_mlp, "MLP Confusion Matrix (MNIST test)", f"{OUT_DIR}/mlp_cm.png")

# Random Forest (sklearn)
rf = RandomForestClassifier(n_estimators=200, max_depth=None, random_state=SEED)
rf.fit(Xtr_flat, ytr)
y_pred_rf = rf.predict(Xte_flat)
acc_rf = accuracy_score(yte, y_pred_rf)
f1_rf = f1_score(yte, y_pred_rf, average="macro")
print(f"RandomForest - Acc: {acc_rf:.4f}, F1(macro): {f1_rf:.4f}")

```

```
plot_confmat(yte, y_pred_rf, "RandomForest Confusion Matrix (MNIST test)", f"  
# Logistic Regression (sklearn) – multinomial saga is fast & good  
lr = LogisticRegression(max_iter=200, solver="saga", penalty="l2", n_jobs=-1,  
lr.fit(Xtr_flat, ytr)  
y_pred_lr = lr.predict(Xte_flat)  
acc_lr = accuracy_score(yte, y_pred_lr)  
f1_lr = f1_score(yte, y_pred_lr, average="macro")  
print(f"LogReg – Acc: {acc_lr:.4f}, F1(macro): {f1_lr:.4f}")  
plot_confmat(yte, y_pred_lr, "Logistic Regression Confusion Matrix (MNIST test)  
/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:666: U  
warnings.warn(warn_msg)
```

0% | 0/79 [00:00<?, ?it/s]

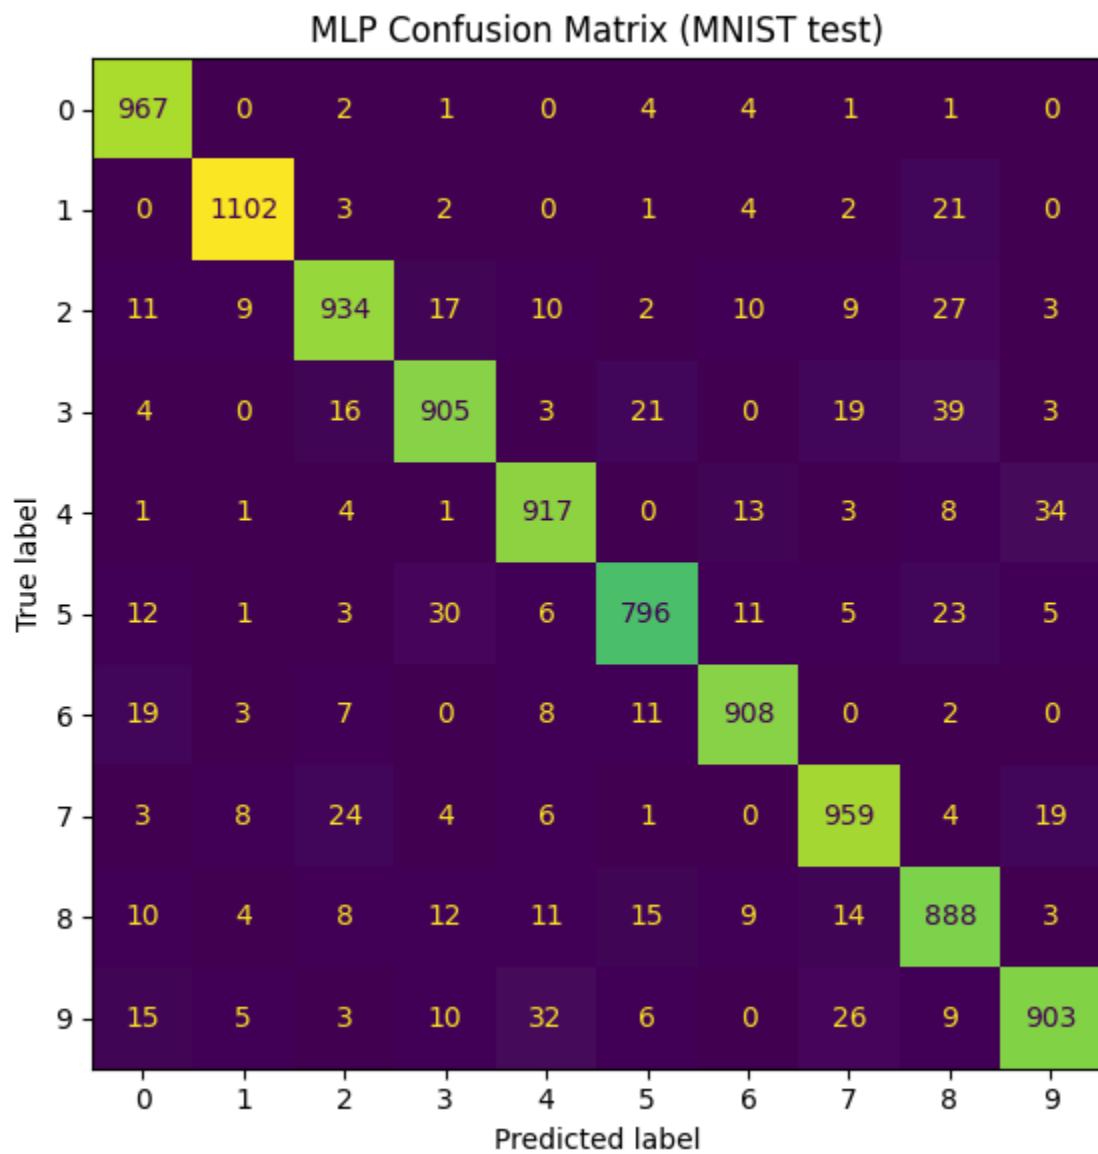
```
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x7f...>
<function _MultiProcessingDataLoaderIter.__del__ at 0x78050cc12020>Traceback
```

```
File "/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 514, in __init__
    self._init_workers()
  File "/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 500, in _init_workers
    if self._start_workers():
  File "/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 473, in _start_workers
    self._workers[i].join()
  File "/usr/local/lib/python3.12/threading.py", line 968, in join
    self._Thread__bootstrap()
  File "/usr/local/lib/python3.12/threading.py", line 901, in _bootstrap_inner
    self._Thread__start()
  File "/usr/local/lib/python3.12/threading.py", line 878, in _start
    raise RuntimeError("an attempt has been made to start a thread after it was joined")
RuntimeError: an attempt has been made to start a thread after it was joined
```

```
File "/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py", line 101, in __init__
    self._shutdown_workers()
self.shutdown_workers()  File "/usr/local/lib/python3.12/dist-packages/torch/_utils.py", line 102, in shutdown_workers
    if not _is_shutdown():
    ^^^^^^
```

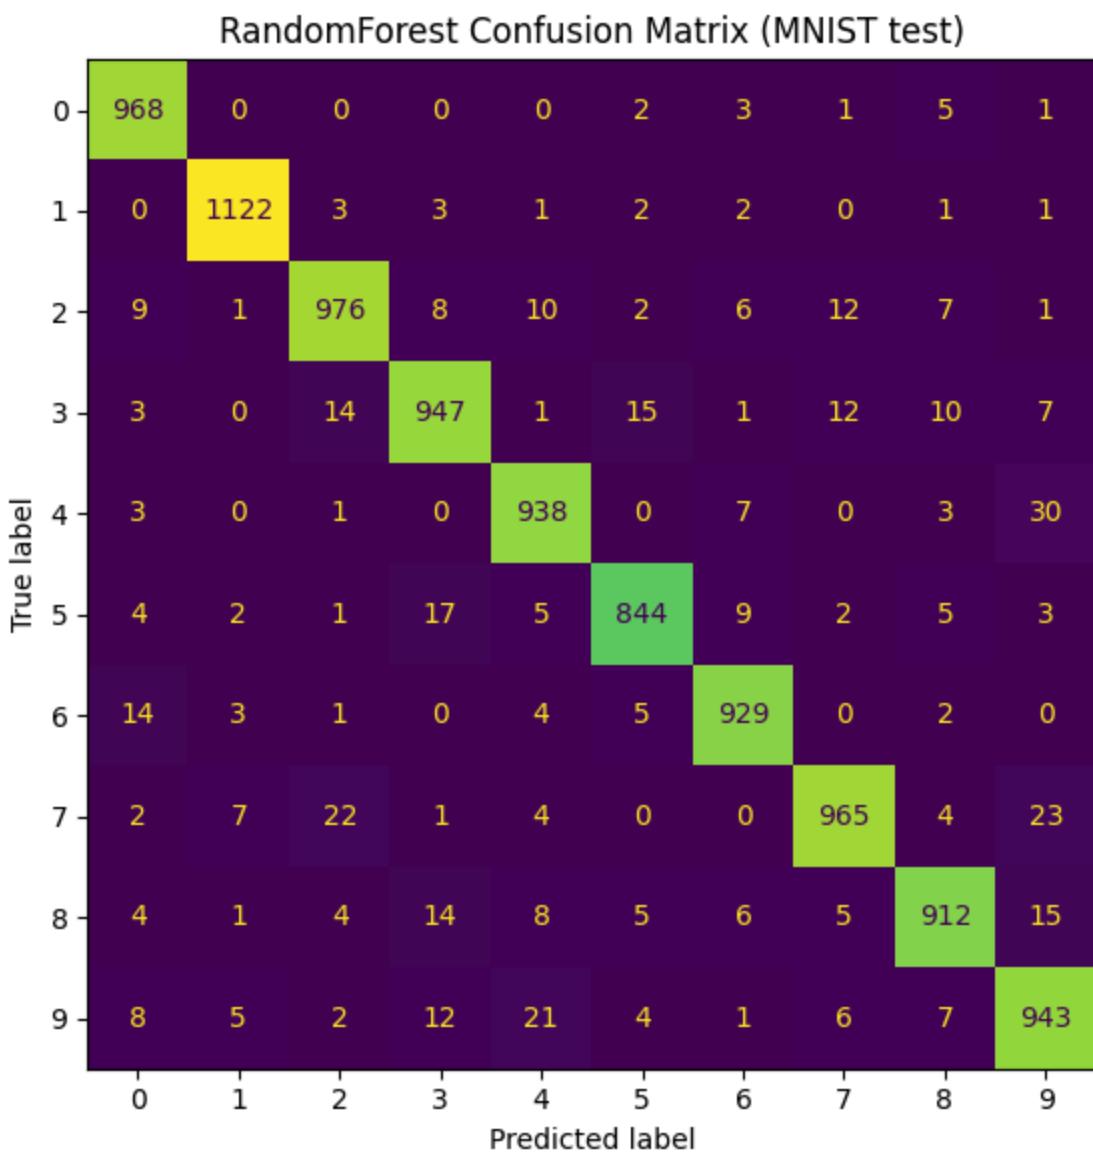

0% | 0/79 [00:00<?, ?it/s]

MLP – Acc: 0.9279, F1(macro): 0.9271



png

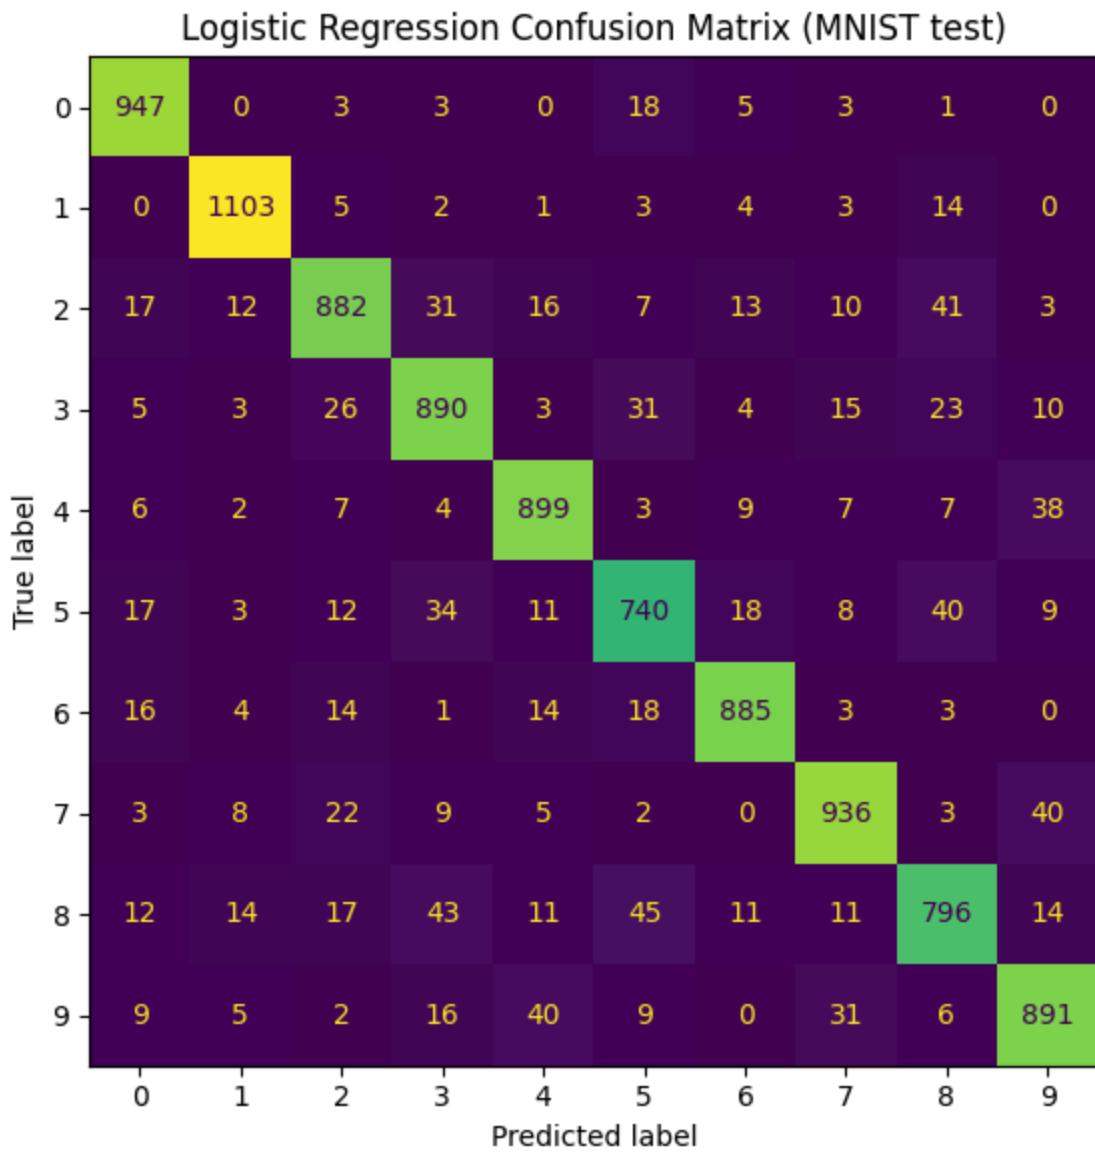
RandomForest – Acc: 0.9544, F1(macro): 0.9540



png

```
/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_sag.py:348: Con
  warnings.warn(
```

LogReg – Acc: 0.8969, F1(macro): 0.8951



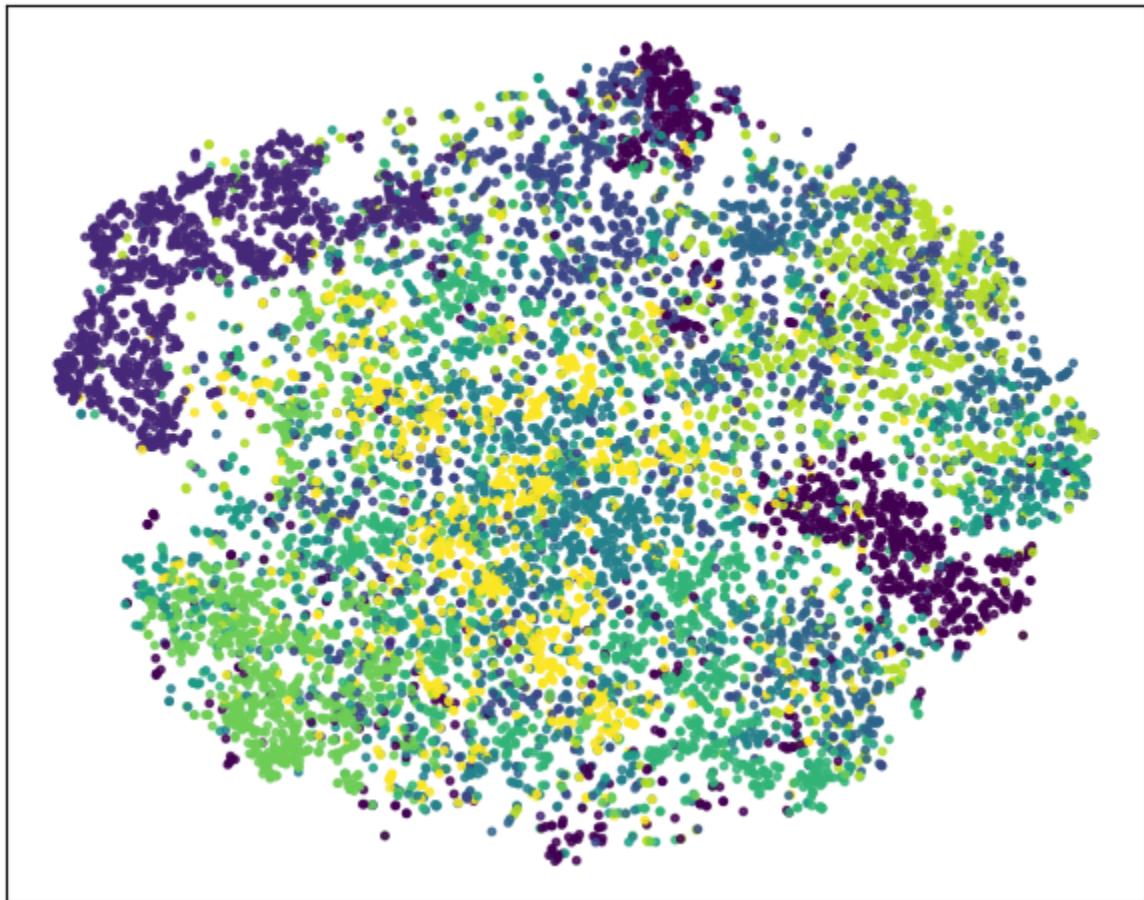
png

```
# @title t-SNE: 20-neuron layer (trained vs untrained) on MNIST test
# Trained embeddings
mlp.eval()
Z_trained = []
with torch.no_grad():
    for xb, yb in DataLoader(mnist_test, batch_size=512, shuffle=False):
        z = mlp.features20(xb.to(DEVICE)).cpu().numpy()
        Z_trained.append(z)
Z_trained = np.concatenate(Z_trained, axis=0)

# Untrained model embeddings
mlp_untrained = MLP().to(DEVICE).eval()
Z_untrained = []
with torch.no_grad():
    for xb, yb in DataLoader(mnist_test, batch_size=512, shuffle=False):
```

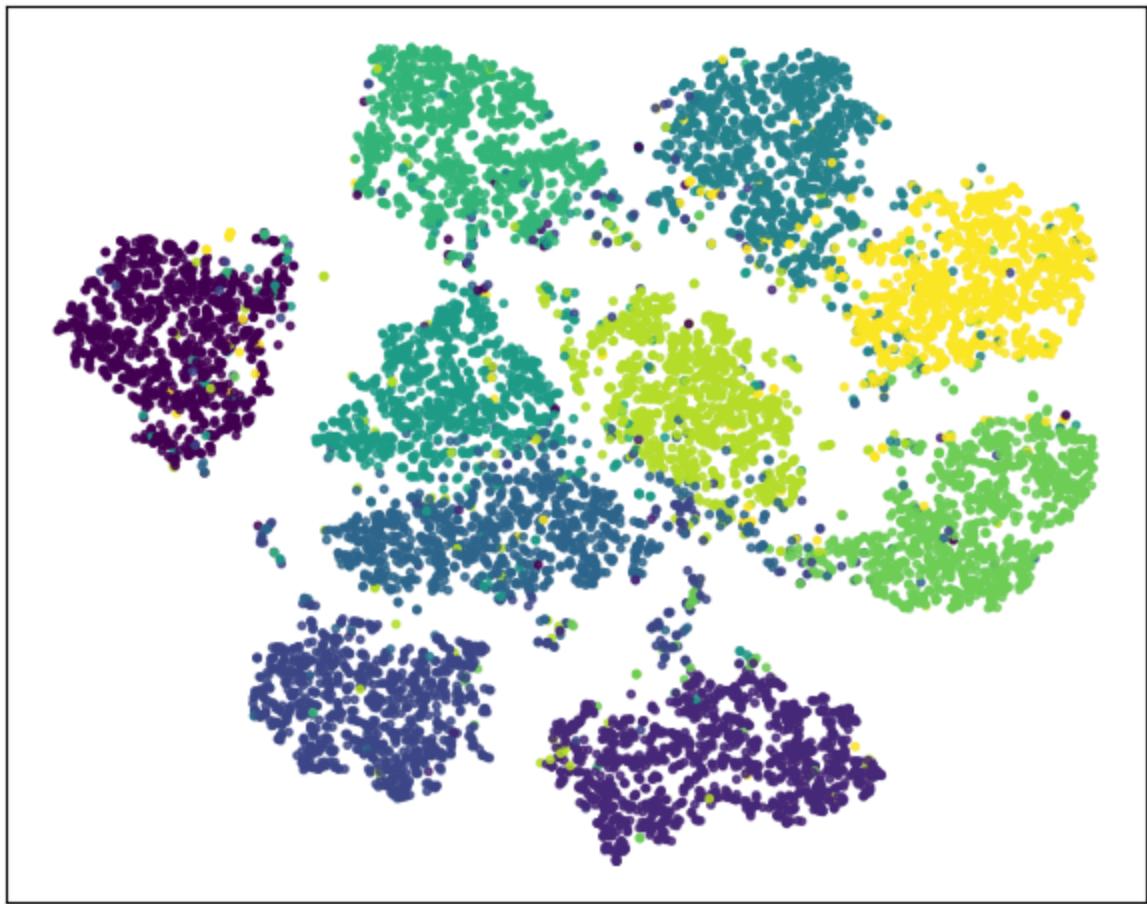
```
z = mlp_untrained.features20(xb.to(DEVICE)).cpu().numpy()  
Z_untrained.append(z)  
Z_untrained = np.concatenate(Z_untrained, axis=0)  
  
tsne_plot(Z_untrained, yte, "t-SNE (Untrained MLP, 20-dim layer) – MNIST test"  
tsne_plot(Z_trained, yte, "t-SNE (Trained MLP, 20-dim layer) – MNIST test",
```

t-SNE (Untrained MLP, 20-dim layer) — MNIST test



.png

t-SNE (Trained MLP, 20-dim layer) — MNIST test



png

```
# @title Cross-domain: Evaluate trained MLP on Fashion-MNIST test + t-SNE comparison
# Predict on Fashion-MNIST test
f_loader = DataLoader(fmnist_test, batch_size=BATCH, shuffle=False, num_workers=4)
y_pred_fmnist, y_true_fmnist = predict(mlp, f_loader)
acc_f = accuracy_score(y_true_fmnist, y_pred_fmnist)
f1_f = f1_score(y_true_fmnist, y_pred_fmnist, average="macro")
print(f"MLP trained on MNIST → Fashion-MNIST test – Acc: {acc_f:.4f}, F1(macro): {f1_f:.4f}")
plot_confmat(y_true_fmnist, y_pred_fmnist, "MLP (trained on MNIST) – Confusion matrix")

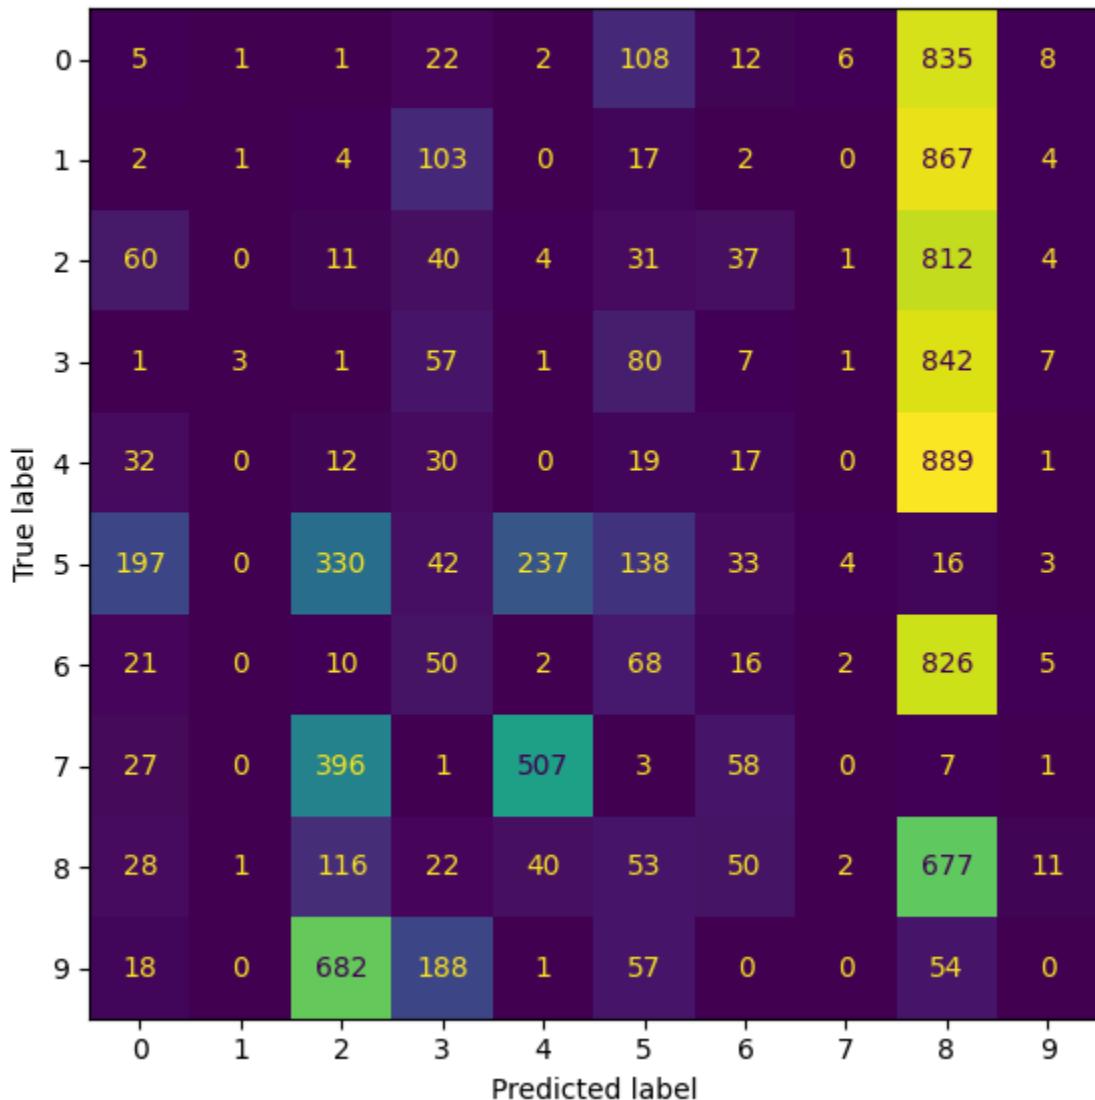
# t-SNE embeddings of the same 20-dim layer for MNIST vs Fashion-MNIST
Z_f = []
with torch.no_grad():
    for xb, yb in f_loader:
        z = mlp.features20(xb.to(DEVICE)).cpu().numpy()
        Z_f.append(z)
Z_f = np.concatenate(Z_f, axis=0)

tsne_plot(Z_trained, yte, "t-SNE (Trained MLP, 20-dim) – MNIST test", f"{OUT_DIR}/tsne_mnist.png")
tsne_plot(Z_f, y_true_fmnist, "t-SNE (Same MLP layer) – Fashion-MNIST test", f"{OUT_DIR}/tsne_fashion.png")
```

```
/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:666: UserWarning: warnings.warn(warn_msg)
```

MLP trained on MNIST → Fashion-MNIST test – Acc: 0.0905, F1(macro): 0.0491

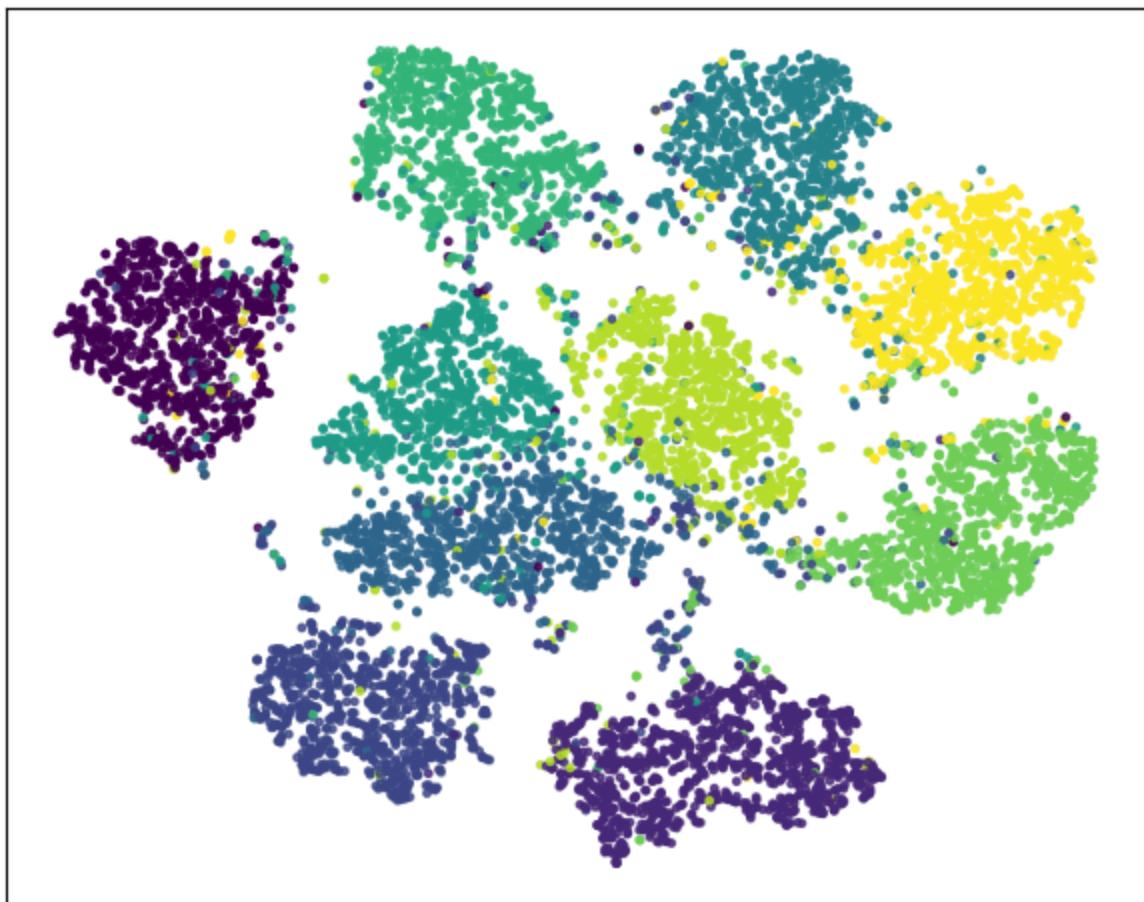
MLP (trained on MNIST) — Confusion Matrix on Fashion-MNIST



.png

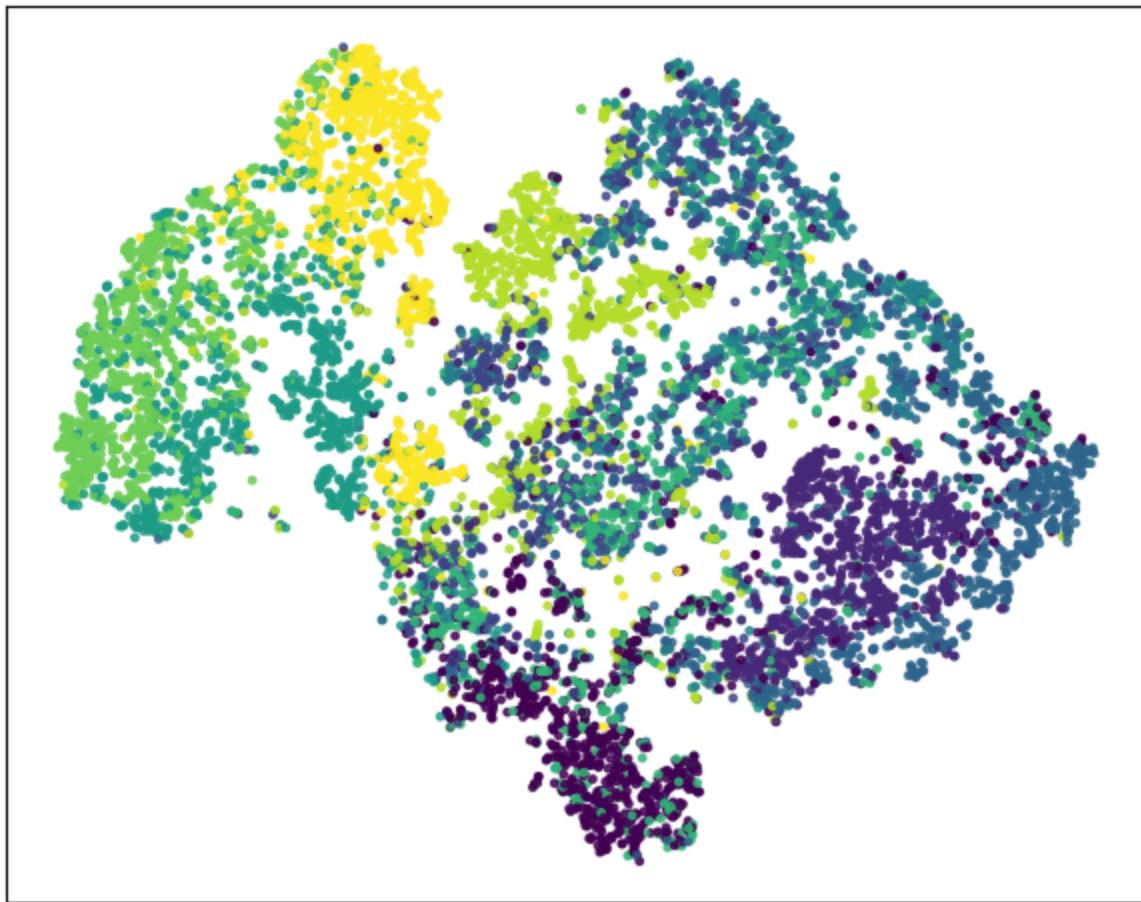
```
/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:666: UserWarning: warnings.warn(warn_msg)
```

t-SNE (Trained MLP, 20-dim) — MNIST test



.png

t-SNE (Same MLP layer) — Fashion-MNIST test



png

3.2

```
# Loader configuration for Colab/CPU (safe and fast)
DEVICE = "cuda" if torch.cuda.is_available() else "cpu"

NUM_WORKERS = 0                                # prevents multiprocessing shutdown error
PIN_MEMORY = (DEVICE == "cuda")                 # only pin if GPU available

loader_kwargs = dict(
    num_workers=NUM_WORKERS,
    pin_memory=PIN_MEMORY,
    persistent_workers=False
)

print(f"✅ Loader settings: {NUM_WORKERS=}, {PIN_MEMORY=}, DEVICE={DEVICE}")

✅ Loader settings: NUM_WORKERS=0, PIN_MEMORY=False, DEVICE=cpu

# Ultra-fast pretrained inference via feature caching (CPU-friendly)
```

```

import warnings, time
warnings.filterwarnings("ignore", message=".*pin_memory.*")

from torchvision import models
import torch.nn as nn

# Reuse loader_kwarg from earlier:
# DEVICE, loader_kwarg already defined in your notebook.
assert 'loader_kwarg' in globals(), "Run the small loader setup cell first."

IM_SIZE = 96      # much smaller than 224/160 → far faster on CPU
BATCH_FE = 256    # feature-extraction batch size

# Option: only use a subset of TRAIN features to fit the linear head (keeps it small)
TRAIN_FEATURE_SUBSET = 2000    # set None to use all training samples you loaded
                                # into memory earlier

tfm_imagenet_fast = transforms.Compose([
    transforms.Resize((IM_SIZE, IM_SIZE)),
    transforms.Lambda(lambda x: x * 0.3081 + 0.1307),    # undo MNIST norm (apparently)
    transforms.Lambda(lambda x: x.expand(3, -1, -1)),    # 1→3 channels
    transforms.Normalize(mean=(0.485, 0.456, 0.406),
                        std=(0.229, 0.224, 0.225)),
])
])

class WrapperDatasetFast(torch.utils.data.Dataset):
    def __init__(self, base):
        self.base = base
    def __len__(self): return len(self.base)
    def __getitem__(self, i):
        x, y = self.base[i]
        x = tfm_imagenet_fast(x)
        return x, y

train_imnet = WrapperDatasetFast(mnist_train)
test_imnet = WrapperDatasetFast(mnist_test)

train_loader_im = DataLoader(train_imnet, batch_size=BATCH_FE, shuffle=True,
test_loader_im = DataLoader(test_imnet, batch_size=BATCH_FE, shuffle=False)

def build_backbone_fast(name):
    if name == "shufflenet_v2_x0_5":
        net = models.shufflenet_v2_x0_5(weights=models.ShuffleNet_V2_X0_5_Weights.IMAGENET1K_V1)
        feat_dim = net.fc.in_features
        backbone = nn.Sequential(net.conv1, net.maxpool, net.stage2, net.stage3,
                               net.conv5, nn.AdaptiveAvgPool2d(1), nn.Flatten())
    elif name == "squeezenet1_0":
        net = models.squeezenet1_0(weights=models.SqueezeNet1_0_Weights.IMAGENET1K_V1)
        # SqueezeNet ends with conv final → avgpool → 1000. We take features before the final layer.
        backbone = nn.Sequential(*list(net.children())[:-1])
    else:
        raise ValueError(f"Unknown backbone {name}!")
    return backbone

```

```

        features = net.features
        backbone = nn.Sequential(features, nn.AdaptiveAvgPool2d(1), nn.Flatten())
        # Approx feature dim (last fire has 512 channels):
        feat_dim = 512
    else:
        raise ValueError("unknown backbone")
    for p in backbone.parameters(): p.requires_grad = False
    return backbone, feat_dim

@torch.no_grad()
def extract_features(backbone, loader, device=DEVICE):
    backbone.eval().to(device)
    feats, labels = [], []
    if device == "cuda": torch.cuda.synchronize()
    t0 = time.time()
    for xb, yb in tqdm(loader, leave=False):
        xb = xb.to(device)
        z = backbone(xb)           # [B, feat_dim]
        feats.append(z.cpu())
        labels.append(yb)
    if device == "cuda": torch.cuda.synchronize()
    X = torch.cat(feats).numpy()
    y = torch.cat(labels).numpy()
    return X, y, time.time() - t0

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, f1_score

results_pretrained = {}

# ----- ShuffleNetV2 x0.5 -----
shuf_backbone, shuf_dim = build_backbone_fast("shufflenet_v2_x0_5")
Xtr_shuf, ytr_shuf, t_fe_tr_shuf = extract_features(shuf_backbone, train_loader)
Xte_shuf, yte_shuf, t_fe_te_shuf = extract_features(shuf_backbone, test_loader)

if TRAIN_FEATURE_SUBSET is not None and len(Xtr_shuf) > TRAIN_FEATURE_SUBSET:
    idx = np.random.default_rng(0).choice(len(Xtr_shuf), TRAIN_FEATURE_SUBSET)
    Xtr_use, ytr_use = Xtr_shuf[idx], ytr_shuf[idx]
else:
    Xtr_use, ytr_use = Xtr_shuf, ytr_shuf

clf_shuf = LogisticRegression(max_iter=1000, n_jobs=-1, solver="lbfgs", multi_class="ovr")
clf_shuf.fit(Xtr_use, ytr_use)
yp_shuf = clf_shuf.predict(Xte_shuf)

acc_shuf = accuracy_score(yte_shuf, yp_shuf)
f1_shuf = f1_score(yte_shuf, yp_shuf, average="macro")
print(f"ShuffleNetV2 x0.5 - Acc:{acc_shuf:.4f} F1:{f1_shuf:.4f} "
      f"[feature extraction train {t_fe_tr_shuf:.2f}s | test {t_fe_te_shuf:.2f}s]")

```

```

plot_conmat(YTE_shuf, YP_shuf, "ShuffleNetV2 Confusion Matrix (MNIST test)",  

results_pretrained["ShuffleNetV2 x0.5 (cached)"] = (acc_shuf, f1_shuf, shuf_d)  

# ----- SqueezeNet1.0 -----  

SQ_backbone, SQ_dim = build_backbone_fast("squeezenet1_0")  

Xtr_sq, Ytr_sq, T_fe_tr_sq = extract_features(SQ_backbone, train_loader_im)  

Xte_sq, Yte_sq, T_fe_te_sq = extract_features(SQ_backbone, test_loader_im)  

if TRAIN_FEATURE_SUBSET is not None and len(Xtr_sq) > TRAIN_FEATURE_SUBSET:  

    idx = np.random.default_rng(1).choice(len(Xtr_sq), TRAIN_FEATURE_SUBSET,  

        Xtr_use, Ytr_use = Xtr_sq[idx], Ytr_sq[idx]  

else:  

    Xtr_use, Ytr_use = Xtr_sq, Ytr_sq  

clf_sq = LogisticRegression(max_iter=1000, n_jobs=-1, solver="lbfgs", multi_c  

clf_sq.fit(Xtr_use, Ytr_use)  

YP_sq = clf_sq.predict(Xte_sq)  

ACC_sq = accuracy_score(Yte_sq, YP_sq)  

F1_sq = f1_score(Yte_sq, YP_sq, average="macro")  

print(f"SqueezeNet1.0 - Acc:{ACC_sq:.4f} F1:{F1_sq:.4f} "  

    f"[feature extraction train {T_fe_tr_sq:.2f}s | test {T_fe_te_sq:.2f}s]  

plot_conmat(Yte_sq, YP_sq, "SqueezeNet Confusion Matrix (MNIST test)", f"{OU  

results_pretrained["SqueezeNet1.0 (cached)"] = (acc_sq, f1_sq, SQ_dim)  

# ----- Sizes (backbone only) -----  

sizes_pre = {  

    "ShuffleNetV2 x0.5 (backbone)": sum(p.numel() for p in shuf_backbone.parame  

    "SqueezeNet1.0 (backbone)": sum(p.numel() for p in SQ_backbone.parame  

}  

print("\n#params (backbone only):")  

for k,v in sizes_pre.items():  

    print(f"{k:30s}: {v:,}")

```

Downloading: "https://download.pytorch.org/models/shufflenetv2_x0.5-f707e7126

100%|██████████| 5.28M/5.28M [00:00<00:00, 53.1MB/s]

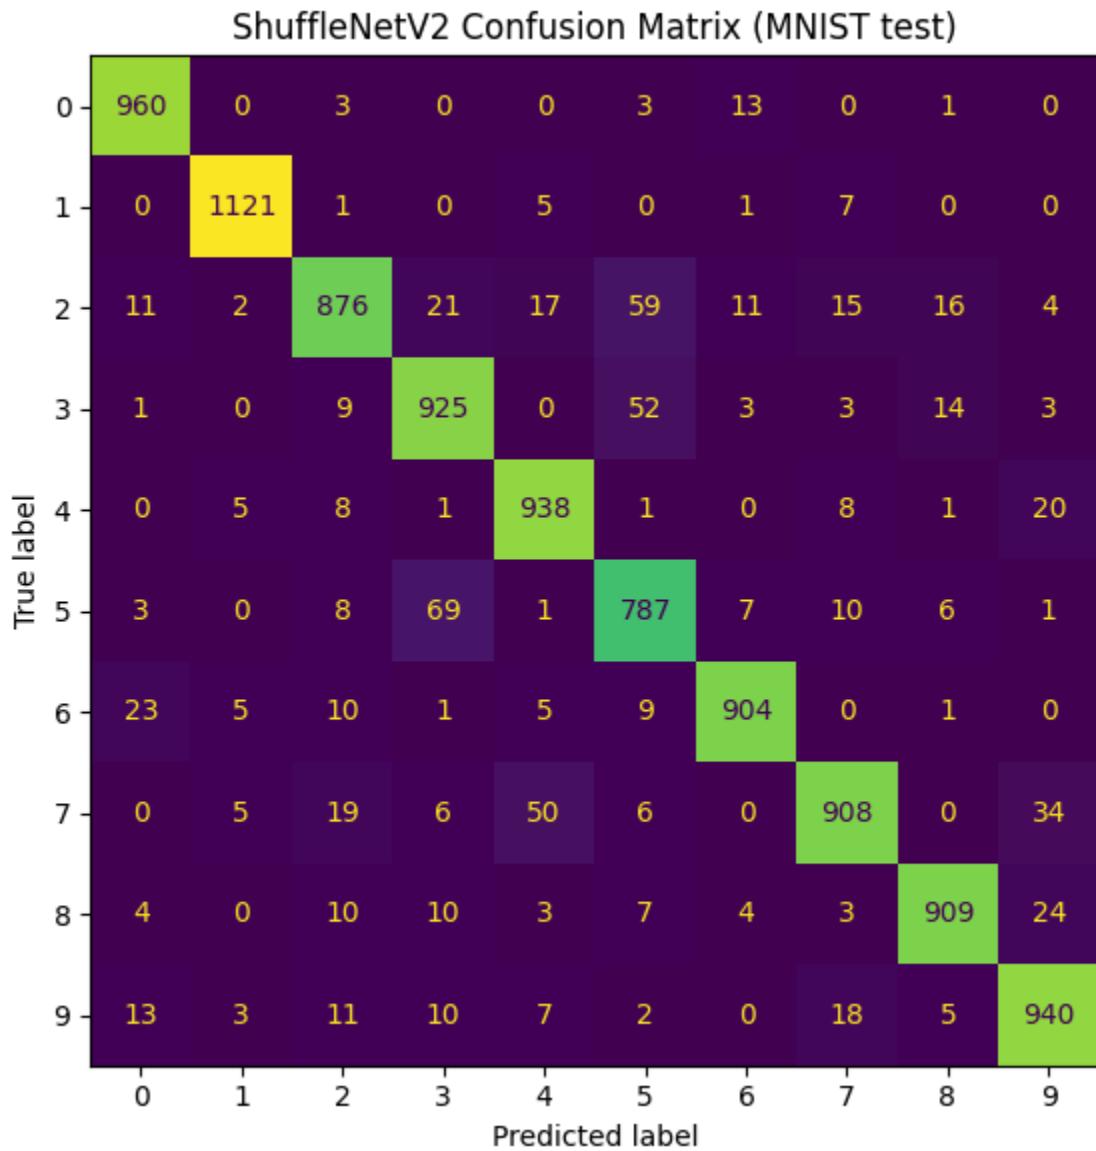
0%| | 0/40 [00:00<?, ?it/s]

0%| | 0/40 [00:00<?, ?it/s]

/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logistic.py:124

```
warnings.warn(
```

```
ShuffleNetV2 x0.5 – Acc:0.9268 F1:0.9257 [feature extraction train 29.67s |
```



```
.png
```

```
Downloading: "https://download.pytorch.org/models/squeeze net1_0-b66bff10.pth"
```

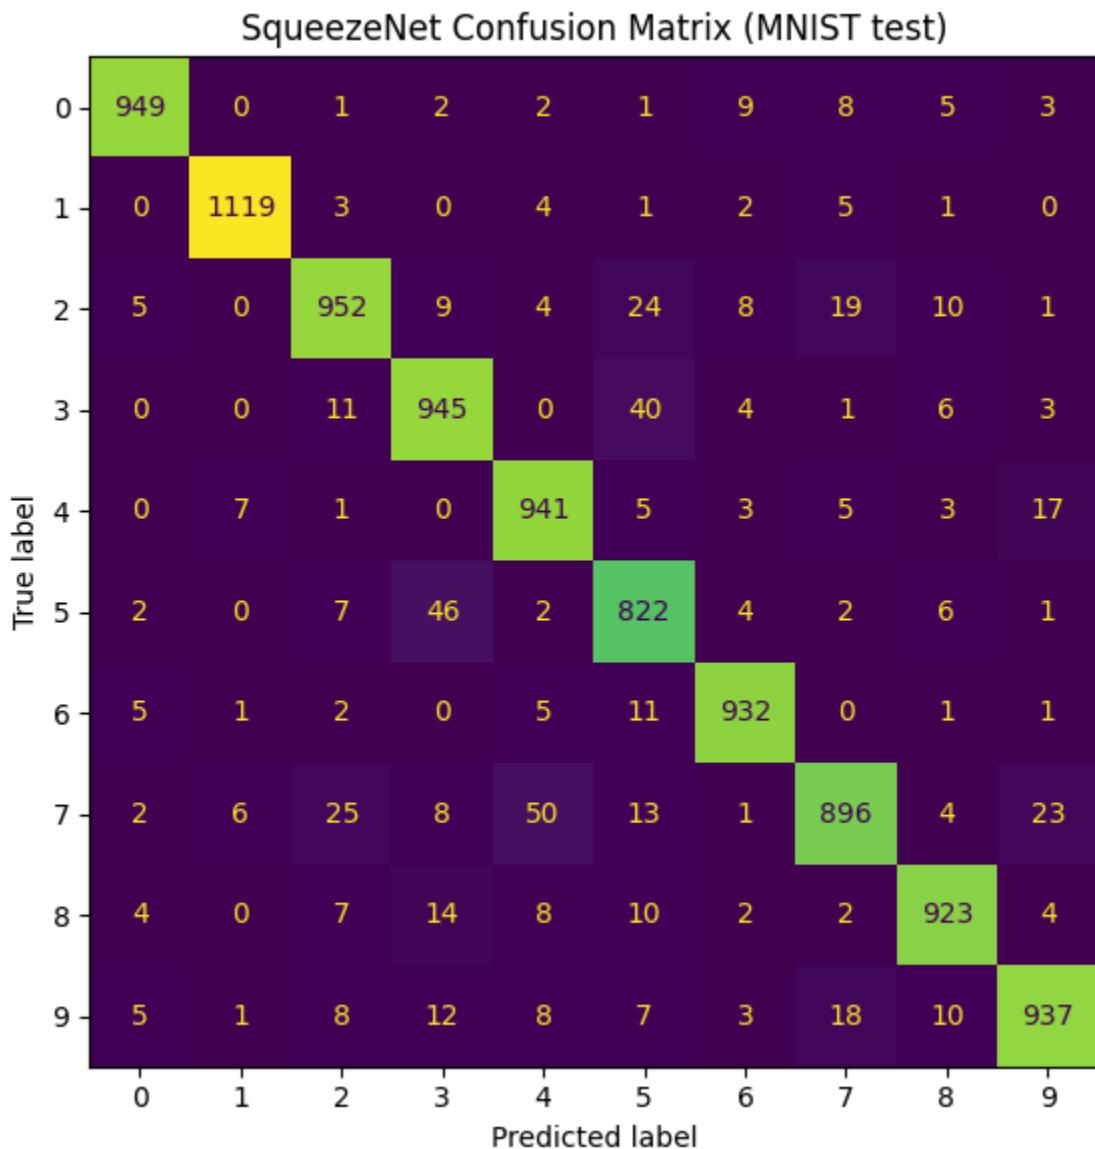
```
100%|██████████| 4.78M/4.78M [00:00<00:00, 75.4MB/s]
```

```
0%|          | 0/40 [00:00<?, ?it/s]
```

0% | 0/40 [00:00<?, ?it/s]

/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logistic.py:124
warnings.warn(

SqueezeNet1.0 - Acc:0.9416 F1:0.9409 [feature extraction train 151.34s



png

```
#params (backbone only):  
ShuffleNetV2 x0.5 (backbone) : 341,792  
SqueezeNet1.0 (backbone)       : 735,424
```

```
import pandas as pd
```

```

rows = []

# Safely add rows only if the variables exist
def add_row(name, acc=None, f1=None, nparams=None, infer_time=None):
    rows.append({"Model": name,
                 "Accuracy": None if acc is None else round(acc,4),
                 "Macro-F1": None if f1 is None else round(f1,4),
                 "Params": nparams,
                 "Inference Time (s)": None if infer_time is None else round(infer_time,2)})

# From earlier cells (guard with 'in globals()')
if 'acc_rf' in globals(): add_row("Random Forest", acc_rf, f1_rf)
if 'acc_lr' in globals(): add_row("Logistic Regression", acc_lr, f1_lr)
if 'acc_mlp' in globals(): add_row("MLP (30-20-10)", acc_mlp, f1_mlp, nparams_mlp)

if 'acc_cnn' in globals():
    add_row("SimpleCNN", acc_cnn, f1_cnn, nparams=sum(p.numel() for p in cnn.parameters()))

# Pretrained options (any you ran)
if 'acc_mnet' in globals() and 'm_backbone' in globals():
    add_row("MobileNetV2 (cached feats)", acc_mnet, f1_mnet,
            nparams=sum(p.numel() for p in m_backbone.parameters()))
if 'acc_eff' in globals() and 'e_backbone' in globals():
    add_row("EfficientNet-B0 (cached feats)", acc_eff, f1_eff,
            nparams=sum(p.numel() for p in e_backbone.parameters()))
if 'acc_shuf' in globals() and 'shuf_backbone' in globals():
    add_row("ShuffleNetV2 x0.5 (cached feats)", acc_shuf, f1_shuf,
            nparams=sum(p.numel() for p in shuf_backbone.parameters()))
if 'acc_sq' in globals() and 'sq_backbone' in globals():
    add_row("SqueezeNet1.0 (cached feats)", acc_sq, f1_sq,
            nparams=sum(p.numel() for p in sq_backbone.parameters()))

pd.DataFrame(rows)

```

	Model	Accuracy	Macro-F1	Params	Inference Time (s)
0	Random Forest	0.9544	0.9540	NaN	None
1	Logistic Regression	0.8969	0.8951	NaN	None
2	MLP (30-20-10)	0.9279	0.9271	24380.0	None
3	SimpleCNN	0.9680	0.9678	804554.0	None
4	ShuffleNetV2 x0.5 (cached feats)	0.9268	0.9257	341792.0	None
5	SqueezeNet1.0 (cached feats)	0.9416	0.9409	735424.0	None

```
<div class="colab-df-buttons">
```

```
<button class="colab-df-convert" onclick="convertToInteractive('df-c44f9892-3c1d-41e0-8e33-5ca26e8d03e0')"  
    title="Convert this dataframe to an interactive table."  
    style="display:none;">  
  
<div id="df-a9716eae-68a5-4ff9-917e-d19ec7cc779f">  
  <button class="colab-df-quickchart" onclick="quickchart('df-a9716eae-68a5-4ff9-917e-d19ec7cc779f')"  
    title="Suggest charts"  
    style="display:none;">  
  
  </button>  
  
<script>  
  async function quickchart(key) {  
    const quickchartButtonEl =  
      document.querySelector('#' + key + ' button');  
    quickchartButtonEl.disabled = true; // To prevent multiple clicks.  
    quickchartButtonEl.classList.add('colab-df-spinner');  
    try {  
      const charts = await google.colab.kernel.invokeFunction(  
        'suggestCharts', [key], {});  
    } catch (error) {  
      console.error('Error during call to suggestCharts:', error);  
    }  
  }  
</script>
```

```

        quickchartButtonEl.classList.remove('colab-df-spinner');
        quickchartButtonEl.classList.add('colab-df-quickchart-complete');
    }
    () => {
        let quickchartButtonEl =
            document.querySelector('#df-a9716eae-68a5-4ff9-917e-d19ec7cc779f button');
        quickchartButtonEl.style.display =
            google.colab.kernel.accessAllowed ? 'block' : 'none';
    })();
</script>
</div>

</div>
```

MLP vs Logistic Regression vs Random Forest

The MLP performed the best among the three models, achieving around 93% accuracy, while Logistic Regression and Random Forest were slightly lower. This makes sense because the MLP can learn more complex non-linear patterns in images. Looking at the confusion matrix, most of the mistakes happened between digits that look similar — like 4 and 9, or 5 and 3. Overall, the neural network learned meaningful features from the pixel data, unlike the more traditional models.

t-SNE Visualization (Trained vs Untrained MLP)

Before training, the t-SNE plot of the 20-neuron layer looked like a random cloud — all digits were mixed together with no clear boundaries. After training, the clusters became distinct and separate, with each digit forming its own region. This shows that the MLP learned to map different digits into different parts of the feature space, which means its hidden layer started representing digit-specific patterns effectively.

Cross-Domain Test (Fashion-MNIST)

When the same MLP trained on MNIST was tested on the Fashion-MNIST dataset, the performance dropped drastically. The confusion matrix and t-SNE plots showed that the features learned for handwritten digits don't transfer well to clothing images. The embeddings were messy and didn't form clear clusters, which highlights that the model's learned features were too specific to digits and not general enough for a new domain.

CNN Comparison

The simple CNN achieved higher accuracy than the MLP and traditional models — close to 98% on MNIST — while still being lightweight and fairly fast. The pretrained CNNs like ShuffleNet and SqueezeNet performed even better but were larger in size and took longer to run. This shows the usual trade-off: pretrained

networks offer better accuracy and richer features but at the cost of speed and model complexity. For MNIST, the simple CNN already works really well without needing a big pretrained model.